

COMMAND AI HUB

BLUEPRINT v1.3 — AI PERSISTENT MEMORY LAYER · ONEDRIVE SYNC

Version 1.3 · 31 May 2026

Build AFTER v1.2.1 passes all 11 checkpoints

Each AI remembers. The team never loses context.

WHY THIS EXISTS: Every AI in the Hub works hard — researches, plans, builds, validates. But today, when a session ends, that knowledge disappears. Blueprint v1.3 fixes that. Each AI gets their own persistent memory file stored in OneDrive. When they reconnect — even months later — they read their file and pick up exactly where they left off. **The team never loses context. Every session makes them smarter.**

PREREQUISITE

This blueprint is built ONLY after Blueprint v1.2.1 passes all 11 checkpoints and is live in Nexus Command. OneDrive connector is already authorised on Kieran's account. Do not begin if v1.2.1 is not yet validated.

THREE MEMORY LAYERS — What Gets Stored and Where

Memory is organised into 3 layers — each serving a different purpose. All stored in OneDrive under a single /5S-Portal/ folder.

Layer	Name	OneDrive Path	What It Contains	Updated When
1	Session Logs	/5S-Portal/AI-Hub/Sessions/YYYY-MM-DD_[session_id].json	Full record of each Hub session — every message, every response, every session end, every deployment.	Automatically updated every time a session ends or a deployment occurs.
2	Role Memory Files	/5S-Portal/AI-Hub/Memory/[ROLE]_memory.md e.g. ENGINEER_memory.md RESEARCHER_memory.md	Each AI's personal memory — what they built, what they learned, what they discovered, what they planned.	Updated whenever an AI writes to their memory file.
3	Blueprint Archive	/5S-Portal/AI-Hub/Blueprints/Blueprint_v[X.X]_[date].pdf	Every blueprint version — the shared source of truth for the whole team.	Updated whenever a new blueprint is finalized and approved.

ONEDRIVE FOLDER STRUCTURE — Exact Layout to Create

Folder / File	Type	Purpose
/5S-Portal/	Root folder	Everything for the portal and AI Hub lives here
/5S-Portal/AI-Hub/	Sub-folder	All AI Hub memory and logs
/5S-Portal/AI-Hub/Sessions/	Sub-folder	One JSON file per session. Auto-created by syncToOneDrive.
/5S-Portal/AI-Hub/Memory/	Sub-folder	One .md file per AI role. The heart of persistent memory.
/5S-Portal/AI-Hub/Memory/ORCHESTRATOR_memory.md	File	ORCHESTRATOR's personal memory log
/5S-Portal/AI-Hub/Memory/RESEARCHER_memory.md	File	RESEARCHER's personal memory log
/5S-Portal/AI-Hub/Memory/ANALYST_memory.md	File	ANALYST's personal memory log
/5S-Portal/AI-Hub/Memory/STRATEGIST_memory.md	File	STRATEGIST's personal memory log
/5S-Portal/AI-Hub/Memory/THINK_TANK_memory.md	File	THINK TANK's personal memory log
/5S-Portal/AI-Hub/Memory/ENGINEER_memory.md	File	ENGINEER's personal memory log
/5S-Portal/AI-Hub/Memory/ARCHITECT_memory.md	File	ARCHITECT's personal memory log
/5S-Portal/AI-Hub/Memory/VALIDATOR_memory.md	File	VALIDATOR's personal memory log
/5S-Portal/AI-Hub/Blueprints/	Sub-folder	Blueprint PDFs — every version archived here

/5S-Portal/Documents/	Sub-folder	Business docs, agreements, compliance files — existing use
-----------------------	------------	--

ROLE MEMORY FILE FORMAT — What Each AI's Memory Contains

Each memory file is plain Markdown — readable by both humans and AI. When an AI is called for the first time in a session, the system loads their memory file and passes the last 500 words as context. They know who they are and what they have done.

```
# ENGINEER_memory.md ## Role Duty: ENGINEER · Model: GPT-5.4 · Provider: OpenAI Lane: Code generation, backend functions, API wiring, structured output. Handoff to: STRATEGIST (for logic review) → VALIDATOR (for final gate) ## What I Have Built - aiCommandCentre function — deployed 31 May 2026. Handles role-aware routing. - processSChatInstruction function — deployed 30 May 2026. S-Chat backend loop. - syncToOneDrive function — deployed [date]. Writes session logs + memory files. ## Patterns That Work - Always use base44.asServiceRole for entity writes in backend functions. - Graph API calls return 202 with empty body — use text() not json(). - AIConnector records drive all role context — read them first before routing. ## Lessons Learned - 31 May 2026: App ID mismatch between sandbox and production caused silent write failures. Fix: always confirm app_id before any write. Production = 69edd16e877d6e4391ad74bd. ## Last 3 Sessions - 31 May 2026 · Built role-aware chat · VALIDATOR: all clear - 30 May 2026 · Fixed S-Chat endpoint · VALIDATOR: all clear - 29 May 2026 · Invoice PDF generator · VALIDATOR: 1 warning resolved
```

TWO NEW FUNCTIONS — What to Build

Function	Trigger	What It Does	Key Steps
syncToOneDrive	Called automatically after S-Chat session logs are saved to OneDrive	Saves the session log as JSON to OneDrive. Also saves the latest S-ChatMessage as a debrief to S-ChatMessages	1. Read S-ChatMessages from S-ChatMessages 2. Write JSON to /AI-Hub/Sessions/[date]_[id].json via Graph API 3. Read ANALYST's debrief from latest SChatMessage 4. Append to each role's _memory.md in /AI-Hub/Memory/ 5. Upload latest blueprint PDF to /AI-Hub/Blueprints/ 6. Return { synced: true, files_written: N }
loadMemory	Called by aiCommandCentre when an AI is first addressed	Reads the AI's memory file from OneDrive, extracts the last 500 words of context, and injects it into the system prompt	1. Read the AI's memory file from OneDrive (e.g. /AI-Hub/Memory/[DUTY_NAME]_memory.md) 2. Fetch /AI-Hub/Memory/[DUTY_NAME]_memory.md from OneDrive via Graph API 3. Extract last 500 words 4. Return { memory_context: string, last_updated: timestamp } 5. aiCommandCentre prepends this to the system prompt for that AI's role

UI ADDITIONS — Small Changes to CommandAIHub.jsx

Only 3 small UI additions needed. No structural changes to the existing layout.

Addition	Where	Detail
Memory indicator on AI status cards	Left panel — each AI card	Small OneDrive icon + 'Memory loaded' (green) or 'No memory yet' (grey). Shows whether memory is loaded
[SYNC TO ONEDRIVE] button	Right panel — below Deploy Log	Azure bg. Fires syncToOneDrive manually. Shows last sync timestamp. Auto-fires after every sync
Memory context banner	Centre panel — above chat	Shows last time an ENGINEER memory loaded — last active [date]. Dismissible. Confirms context is loaded

RECONNECT SCENARIO — What Happens When an AI Comes Back

This is the scenario Kieran asked about. An AI gets disconnected. Comes back months later. Here is exactly what happens.

Step	What Happens	Who Does It
1	Kieran types 'ENGINEER —' in the Hub chat	Kieran
2	aiCommandCentre detects ENGINEER is being called. Checks if ENGINEER has a memory file in OneDrive	aiCommandCentre
3	loadMemory fetches ENGINEER_memory.md from OneDrive. Extracts last 500 words of context	loadMemory function
4	Memory context is prepended to ENGINEER's system prompt before they respond.	aiCommandCentre function

5	ENGINEER responds — in character, with full awareness of past work: "ENGINEER here. Back online. I can see we've built aiCommandCentre, processSChatInstruction, and syncToOneDrive. Last session was [date]"	ENGINEER
6	Blue memory banner appears in chat: 'ENGINEER memory loaded — last active [date]'	CommandAIHub.jsx
7	Session continues. ENGINEER works as if they never left.	Full team

BUILDER PREFIX — Paste This to Implement Blueprint v1.3

"Build Blueprint v1.3 — AI Persistent Memory Layer for Command AI Hub (Nexus Command app). PREREQUISITE: Blueprint v1.2.1 must be live and all 11 checkpoints passed before starting this. 1. CREATE ONEDRIVE FOLDER STRUCTURE via Microsoft Graph API: /5S-Portal/AI-Hub/Sessions/ · /5S-Portal/AI-Hub/Memory/ · /5S-Portal/AI-Hub/Blueprints/ Create one [ROLE]_memory.md file per AI (8 files) with role header, lane, and handoff_to pre-filled. 2. BUILD syncToOneDrive FUNCTION: Triggered after every deploy. Reads SChatMessage entity for session. Writes JSON session log to Sessions folder. Appends ANALYST debrief to relevant role memory.md files. Uploads latest blueprint to Blueprints folder. 3. BUILD loadMemory FUNCTION: Called by aiCommandCentre when an AI is first addressed. Fetches [ROLE]_memory.md from OneDrive. Returns last 500 words as memory_context string. aiCommandCentre prepends to system prompt. 4. UPDATE aiCommandCentre FUNCTION: On first call to any duty name in a session — call loadMemory, inject context into system prompt. 5. UPDATE CommandAIHub.jsx — 3 small additions only: (a) Memory indicator on each AI status card in left panel — OneDrive icon + loaded/not loaded state. (b) [SYNC TO ONEDRIVE] button in right panel below Deploy Log — fires syncToOneDrive manually. (c) Memory context banner in centre panel — slim blue bar on first AI response confirming memory loaded. Use existing OneDrive connector (already authorised). Same design system: #e8e6fe, #5e50fb, #0078d4 for OneDrive elements. Exo 2 headlines, Montserrat body. No emoji."

WHAT TO TEST — After Builder Delivers v1.3

Test	What Kieran Does	What Should Happen	Pass Condition
T1 — Folder created	Open OneDrive after build	All 3 folders exist. 8 memory files exist in /Memory/	Files visible in OneDrive ■
T2 — Memory loaded	Type 'ENGINEER —' in Hub chat	Blue memory banner appears. ENGINEER mentions past work in response	Engineer responds in context ■
T3 — Session sync	Fire a test deploy, then check OneDrive	JSON file appears in /Sessions/. Memory files updated	Files written to OneDrive within 30 seconds ■
T4 — Reconnect	Clear browser cache, reopen Hub	ENGINEER loads Memory from OneDrive fresh, responds with full history	Memory fully restored ■
T5 — VALIDATOR	Run VALIDATOR after build	All 4 Gatekeeper Panel checks green. TestLog 4 records	VALIDATOR returns ALL CLEAR ■

BUILD ROADMAP — The Full Journey

Blueprint	Focus	Status	Prerequisite
v1.2.1	Role-aware chat + Gatekeeper Panel + AI Family voices	Send to builder now	—
v1.3	AI Persistent Memory — OneDrive sync + role memory files document — build after v1.2.1	11.2.1 all 11 checkpoints passed	
v1.4	DESIGNER (Canva AI) + BRAND DIRECTOR (Looka AI) + Branding	Deployed when v1.3 is live	v1.3 all 11 checkpoints passed
v1.5	Cross-portal sync — Hub memories feed into 5S Portal sessions	Decision	v1.4 live