

NEXUS COMMAND

AI MULTI-AGENT BLUEPRINT v1.5 — LOCKED

All 5 decisions confirmed by Kieran · 1 June 2026 · SIMPLEX-ITY Confidential

SECTION 1 — KIERAN'S 5 DECISIONS (ALL LOCKED)

#	Topic	Kieran's Decision	Simpee's Action
Q1	QC AI — Tool vs AI	Find a real AI (not just a static tool) that can join the team and review code collaboratively.	■ RECOMMENDATION: CodeRabbit AI — true AI reviewer, free for open-source, \$24/mo for private. Acts as a team member, leaves comments, learns codebase. Invite as QC INSPECTOR.
Q2	Logger / Memory AI	Confirm Gemini joins as Logger/Memory AI — but want to know if paid or free.	■ Gemini API has a FREE tier (15 RPM, 1M tokens/day). For our volume: FREE. Gemini 2.0 Flash is free for development. Assign dual role: RESEARCHER + LOGGER/MEMORY. No extra cost.
Q3	Ethics AI Scope	General compliance handled by Research team. Ethics AI is optional/future — not a priority now.	■ DEFERRED: Remove Ethics AI as a mandatory gate in v1.5. Research team (Gemini + Claude Sonnet) handles compliance checks when required. Reintroduce in v1.7 if needed.
Q4	Monitor AI Conflict	Agreed — Copilot = Validator only. Claude Sonnet = Monitor AI.	■ LOCKED: Claude Sonnet gets dual role — ANALYST + MONITOR AI. AIConnector updated. Copilot = VALIDATOR only. No conflicts.
Q5	Feedback Loop Exit	ALL AIs must participate in review before escalating to Kieran. Give extra solution. Then WhatsApp with final yes/no + reason + alternatives.	■ LOCKED: 3-failure rule upgraded to FULL TEAM REVIEW. See escalation protocol below.

SECTION 2 — QC AI RECOMMENDATION (Answering Q1)

Kieran asked: 'Find a real AI that can join the team, not just a static tool.' Here are the options:

Tool	What it does	Cost	Verdict
CodeRabbit AI	AI code reviewer that reads your full codebase, understands context, leaves inline comments like a real team member.	Free (open source) / \$24/mo (private)	■ RECOMMENDED — hire as QC INSPECTOR
Greptile	Full-codebase context AI reviewer, self-hostable, strong for small teams.	Free tier available	■ BACKUP if CodeRabbit not suitable
Qodo (formerly CodiumAI)	AI that generates tests + reviews. Strong on test coverage.	\$19/mo	■ OPTIONAL — add as TEST WRITER role
Snyk Code	Security-focused AI scanner. Finds vulnerabilities, not general code quality.	Free tier (limited)	■ Add to Ethics/Security layer later

■ SIMPEE RECOMMENDATION: Start with CodeRabbit AI — it's the closest to a true AI team member for code review. It reads your whole codebase, understands context, posts comments like a colleague, and learns over time. Cost: FREE for open-source repos, \$24/mo for private. Action: Go to coderabbit.ai and connect your Nexus Command repo.

SECTION 3 — ESCALATION PROTOCOL (Answering Q5 — FULL TEAM REVIEW)

Rule: 3 consecutive failures → ALL 8 AIs participate in review before Simpee escalates to Kieran. Here is the exact 9-step protocol:

STEP 1 — ENGINEER self-diagnoses	GPT-5.4 reviews own code, identifies root cause, proposes a fix. Logs in TasiLog
STEP 2 — ARCHITECT reviews	
STEP 3 — ANALYST assesses	Claude Sonnet (now also MONITOR AI) reviews error pattern. Is this a 1-off bug or a recurring systemic issue?
STEP 4 — RESEARCHER investigates	
STEP 5 — THINK TANK brainstorms	Claude Opus explores unconventional approaches. What if we scrap this approach entirely? Proposes 1 alternative path
STEP 6 — STRATEGIST evaluates	
STEP 7 — VALIDATOR final check	Copilot reviews all proposed solutions for syntax validity, security, Azure readiness. Marks each SAFE / RISKY
STEP 8 — ORCHESTRATOR consolidates	
STEP 9 — WhatsApp to Kieran	Simpee sends Kieran a full brief. ■ ISSUE: [what failed + reason] ■ TEAM REVIEWED: 7 AIs participated ■ SOLUTION A (recommended): [description] — safety: SAFE ■ SOLUTION B (alternative): [description] — safety: SAFE ■■ SOLUTION C (risky): [description] — safety: RISKY ■ YOUR CALL: Approve A / Approve B / Abort this feature?

The WhatsApp to Kieran is NOT a panic message. It is a full briefing with all options ranked, safety-checked, and ready for his one-word decision. Kieran never troubleshoots — he just decides.

SECTION 4 — FINAL TEAM ROSTER v1.5 LOCKED

Duty Name	Model	Status	Role
ORCHESTRATOR	Simpee (Base44)	■ ACTIVE	CEO — governance, routing, deploy gating, WhatsApp escalation
RESEARCHER	Gemini 2.0 Flash	■ ACTIVE + DUAL	External data, benchmarks + Logger/Memory AI (free tier)
ANALYST	Claude Sonnet 4.6	■ ACTIVE + DUAL	Feasibility, risk assessment + Monitor AI (cross-role watch)
STRATEGIST	Claude Opus 4.6	■ ACTIVE	ROI framing, opportunity cost, solution ranking
THINK TANK	Claude Opus 4.8	■ ACTIVE	Deep reasoning, unconventional approaches
ENGINEER	GPT-5.4	■ ACTIVE	Code writing, debugging, implementation
ARCHITECT	GPT-5.5	■ ACTIVE	System design, data architecture, infrastructure
VALIDATOR	Copilot / Edge	■ ACTIVE	Syntax, logic, security, Azure readiness — FINAL GATE
QC INSPECTOR	CodeRabbit AI	■ HIRE NOW — \$24/mo	AI code review, inline comments, codebase-aware QC
ETHICS AI	Research Team	■ DEFERRED to v1.7	Handled by Gemini + Claude Sonnet when required

SECTION 5 — NEXT STEPS TO v1.6

#	When	Action	Owner
1	TODAY	Kieran registers CodeRabbit AI at coderabbit.ai — connect to Nexus Command GitHub repo	■ Kieran action

#	When	Action	Owner
2	TODAY	Simpee updates AIConnector entity: Gemini = dual (Researcher + Logger), Claude Sonnet = dual (Analyst + Monitor)	■ Simpee action
3	THIS WEEK	Simpee creates KnowledgeLog entity for Role 6 unified truth layer	■ Simpee action
4	THIS WEEK	Full 9-step escalation protocol tested on a sample failed build in Nexus sandbox	■ Joint
5	THIS WEEK	Simpee upgrades aiCommandCentre function to log all runs to TestLog with team participation fields	■ Simpee action
6	NEXT SPRINT	Run full 8-role dry-run build. All 11 checkpoints must pass.	■ Joint
7	v1.6 TARGET	If dry-run passes → v1.6 PRODUCTION READY. Promote to 5S Portal.	■ Kieran approves

GOVERNANCE RULES v1.5 — FINAL

- RULE 1: No deploy without QC Inspector (CodeRabbit) + Validator (Copilot) + Orchestrator (Simpee) all approving
- RULE 2: 3-failure rule → Full 9-step team review → WhatsApp Kieran with ranked options
- RULE 3: Copilot = Validator ONLY. Claude Sonnet = Analyst + Monitor AI. No conflicts.
- RULE 4: Gemini = Researcher + Logger/Memory AI. Free tier. No extra cost.
- RULE 5: Ethics AI deferred to v1.7. Research team handles compliance checks ad-hoc.
- RULE 6: Kieran never troubleshoots. He receives a briefing and makes one decision.
- RULE 7: No AI approves its own work. Reviewer must be a different role.
- RULE 8: Blueprint-first. All changes require a blueprint before any code is written.
- RULE 9: All escalations are logged to TestLog with all 8 AI participation fields.
- RULE 10: Knowledge AI (Gemini + Claude Sonnet) maintains single truth layer in KnowledgeLog entity.