

NEXUS COMMAND — AI MULTI-AGENT BLUEPRINT v1.5

Generated by Simpee · SIMPLEX-ITY · 1 June 2026 · Confidential

WHAT'S NEW IN v1.5 (vs v1.4)

- QC INSPECTOR role formally defined — SonarQube + DeepSource as external QC layer
- 8 Roles fully mapped with clear workflows, outputs, and assigned models
- Logger AI + Memory AI → assigned to Gemini 2.0 Pro (dual role, pending confirmation)
- Monitor AI → assigned to Copilot dual role (pending Q4 clarification)
- Knowledge AI → Gemini + Claude Sonnet joint session protocol
- Ethics & Compliance AI → placeholder with recommended implementation path
- Performance Optimizer AI → GPT-5.4 + Copilot dual-diagnostics
- Feedback loop break condition added (3-failure rule → human escalation)
- 5 open questions documented for Kieran's decision before v1.5 is locked
- Hiring guide for all placeholder roles included

FULL AI TEAM ROSTER — 15 MEMBERS

Duty Name	Model	Status	Primary Role
ORCHESTRATOR	Simpee (Base44)	■ ACTIVE	CEO — governance, routing, deploy gating
RESEARCHER	Gemini 2.0 Pro	■ ACTIVE	External data, benchmarks, market insights
ANALYST	Claude Sonnet 4.6	■ ACTIVE	Feasibility, complexity, risk assessment
STRATEGIST	Claude Opus 4.6	■ ACTIVE	ROI framing, opportunity cost, dependencies
THINK TANK	Claude Opus 4.8	■ ACTIVE	Deep reasoning, unconventional approaches
ENGINEER	GPT-5.4	■ ACTIVE	Code writing, debugging, implementation
ARCHITECT	GPT-5.5	■ ACTIVE	System design, data architecture, infra
VALIDATOR	Copilot / Edge	■ ACTIVE	Syntax, logic, security, Azure readiness
QC INSPECTOR	SonarQube + DeepSource	■ PLACEHOLDER	Static analysis, bug/security scanning
LOGGER	Gemini 2.0 Pro (dual)	■ TO ASSIGN	Error logging, session memory, audit trail
MEMORY AI	Gemini 2.0 Pro (dual)	■ TO ASSIGN	Long-context retention, lesson recall
MONITOR AI	Copilot (dual)	■ TO ASSIGN	Cross-role error escalation, intervention
KNOWLEDGE AI	Gemini + Claude Sonnet	■ TO ASSIGN	Unified truth layer, insight consolidation
ETHICS AI	TBD — Rules Engine	■ HIRE NEEDED	Bias check, compliance, HK data privacy
PERF OPTIMIZER	GPT-5.4 + Copilot	■ TO ASSIGN	Efficiency monitoring, resource usage

8 OPERATIONAL ROLES

Role 1: Research, Opportunity Analysis & Value Estimation	Flow: RESEARCHER (Gemini) → ANALYST (Claude Sonnet) → STRATEGIST (Claude Opus)
Output:	Structured insights + ROI estimation + feasibility report

Role 2: Engineering, Learning & Validation	Flow: ENGINEER (GPT-5.4) → ARCHITECT (GPT-5.5) → VALIDATOR (Copilot)
Output:	Validated code modules ready for QC inspection
Role 3: QC Inspection & Approval (3-Gate Deploy Rule)	Flow: QC INSPECTOR (SonarQube/DeepSource) → VALIDATOR (Copilot) → ORCHESTRATOR (Simpee)
Output:	No deploy unless ALL THREE approve. Any failure = return to Role 2.
Role 4: Error Logging & Continuous Learning	Flow: LOGGER AI (Gemini) → MEMORY AI (Gemini) → ANALYST (Claude Sonnet)
Output:	Lessons fed back into Role 1 & Role 2. Max 3 retries before human escalation.
Role 5: Problem-Solving Suggestion Team	Flow: MONITOR AI (Copilot) → THINK TANK (Claude Opus) → ANALYST (Claude Sonnet)
Output:	Consolidated corrective summary for process team. Triggered on escalation.
Role 6: Knowledge Integration	Flow: RESEARCHER (Gemini) + ANALYST (Claude Sonnet) → KnowledgeLog entity
Output:	Single unified truth layer. All insights consolidated. No contradictions.
Role 7: Ethics & Compliance Review	Flow: ETHICS AI (Rules Engine) + VALIDATOR (Copilot) → sign-off required
Output:	Bias check, HK PDPO compliance, data privacy. Required before any deploy.
Role 8: Performance Optimization	Flow: ENGINEER (GPT-5.4) + VALIDATOR (Copilot) → performance diagnostics
Output:	Efficiency monitoring, resource usage, response time benchmarking.

FEEDBACK LOOP ARCHITECTURE

- Role 4 (Logger/Memory) → feeds lessons into Role 1 (Research) and Role 2 (Engineering)
- Role 5 (Monitor) → watches ALL roles, intervenes when errors escalate beyond 3 retries
- Role 6 (Knowledge) → ensures unified knowledge base, prevents contradictions across roles
- Role 7 (Ethics) → enforces compliance BEFORE any deployment reaches Role 3 gate
- Role 8 (Performance) → validates efficiency AFTER Role 3 approval, BEFORE final release
- Simpee (Orchestrator) → coordinates all roles, holds veto on any stage, owns final deploy
- BREAK CONDITION: 3 consecutive failures on same task → Simpee pauses loop + WhatsApps Kieran

GOVERNANCE RULES (UPDATED v1.5)

- RULE 1: No deployment without QC Inspector + Validator (Copilot) + Orchestrator (Simpee) all approving
- RULE 2: Ethics AI must sign off before any code reaches the deploy gate
- RULE 3: All errors must be logged to TestLog before next cycle begins
- RULE 4: Problem-Solving Team (Role 5) reviews all escalations, provides corrective path
- RULE 5: Performance Optimizer validates efficiency before final release
- RULE 6: No AI approves its own work — reviewer must be a different role
- RULE 7: Feedback loop max 3 retries — then human escalation to Kieran via WhatsApp
- RULE 8: Blueprint-first — all changes require a blueprint before any code is written
- RULE 9: Knowledge AI maintains single truth layer — contradicting records are flagged
- RULE 10: Copilot / Validator is the FINAL gate — even after QC Inspector passes

HIRING GUIDE — PLACEHOLDER ROLES

These are the roles Kieran needs to 'hire' — either by API integration or role reassignment:

Role Needed	Recommended Model/Tool	How to Hire	Priority
QC INSPECTOR	SonarQube Cloud (free tier) or DeepSource	Register at sonarcloud.io — connect via GitHub. Add webhook to Nexus Command.	■ HIGH
ETHICS AI	OpenAI Moderation API or custom Rules Engine	Use OpenAI /v1/moderations endpoint. Build a rules JSON for HK compliance + data privacy.	■ MEDIUM
LOGGER / MEMORY AI	Gemini 2.0 Pro (dual-role with RESEARCHER)	Already hired — just assign dual duty in AIConnector entity. Add memory_store field.	■ LOW (reassign)
MONITOR AI	Copilot (dual-role with VALIDATOR)	Already hired — add monitor_active=true in AIConnector. Give cross-role alert permission.	■ LOW (reassign)
KNOWLEDGE AI	Gemini + Claude Sonnet (joint)	Already hired — define joint session protocol. Both write to a shared KnowledgeLog entity.	■ LOW (reassign)

OPEN QUESTIONS FOR KIERAN — DECIDE BEFORE v1.5 LOCKS

Simpee needs your answers on these 5 points before the blueprint is finalised and implemented:

Q1	QC AI Integration	SonarQube/DeepSource are external tools, not chat AIs. Do you want API integration (webhook-based), or assign an AI to simulate QC behaviour inside the hub?
Q2	Logger / Memory AI	Recommend Gemini 2.0 Pro takes dual duty (Researcher + Logger/Memory). It handles long-context best. Confirm?
Q3	Ethics AI Scope	Is Ethics AI for: (A) Code ethics — bias in AI outputs, (B) Business ethics — HK compliance, PDPO data privacy, or (C) Both? This shapes the rules engine.
Q4	Monitor AI Conflict	Copilot is currently Validator AND Monitor AI. Double duty risks conflicts. Consider separating: Copilot = Validator only, Claude Sonnet = Monitor AI. Agree?
Q5	Feedback Loop Break Condition	Role 4 → Role 1/2 feedback loop has no defined exit. What's the max retry count before Simpee escalates to Kieran? Suggest: 3 failures = human escalation.

NEXT STEPS — v1.5 → v1.6 ROADMAP

- Kieran answers 5 open questions above
- Simpee updates AIConnector entity with dual-role assignments (Logger, Monitor, Knowledge)
- Kieran registers SonarCloud (free) and connects to Nexus Command repo
- Simpee builds Ethics Rules Engine as a backend function (simple JSON rules + OpenAI Moderation API)
- Simpee creates KnowledgeLog entity for Role 6 unified truth layer
- Full 8-role test run on a sample build task in sandbox
- If all 11 checkpoints pass → promote to v1.6 PRODUCTION READY